

COVID-19 Diagnosis

Through Images and Audio

Chris Lewis

Some COVID Statistics

100,000,000+
Confirmed cases worldwide

2,000,000+
Deaths worldwide

20,000+
Daily ICU Hospitalizations



In the United States since early
December 2020

100,000+
Daily Hospitalizations

In the United States since early
December 2020

What is COVID-19?

- A strain of a novel coronavirus that has not been previously detected in humans
- Easily transmissible
- Highly contagious
- Not all symptoms are present in those who become infected
- Significant percentage of those infected are asymptomatic




Modes of Transmission:

1. Contact (physical contact or contaminated surfaces)
2. Droplet
3. Airborne

Severity and Symptoms:

- Range from Asymptomatic → Severe
- Very few telltale symptoms:
 - Immediate and significant loss of taste and / or smell
 - “Covid toes”
 - Distinctive Dry cough and shortness of breath



Types of Testing for COVID



Viral Testing:


- Used if suspected of having a current COVID infection
- Two subtypes:
 - Molecular (PCR)
 - Antigen

Antibody Testing:

- Used to determine if a past COVID infection occurred

X-Ray Image Classification via Computer Vision

1. Gathered dataset via Kaggle's API
2. Contained 3800+ high quality chest x-ray images
3. Multiple classes: Healthy, Viral Pneumonia, COVID
4. Used CLAHE as a preprocessing technique
5. Created a Sequential Convolutional Neural Network




Ex. of CNN
architecture



X-ray Model Evaluation Results

Classification Report					
	precision	recall	f1-score	support	
0	0.98	0.98	0.98	240	
1	0.95	0.98	0.96	269	
2	0.98	0.94	0.96	270	
accuracy			0.97	779	
macro avg	0.97	0.97	0.97	779	
weighted avg	0.97	0.97	0.97	779	



- 97% Accuracy
- 98% Recall rate on COVID class
- Less than 2% False Negatives for COVID class

RECOMMENDATION:

- Using Computer Vision with x-rays images is a great way to diagnose patients for many diseases.
- Our model is able to achieve an accuracy equal to a professional when diagnosing chest x-rays for our three classes.

Things to Consider:



Implications of getting a chest X-ray if you are COVID positive:

- Exposure to those in the building
- Getting an x-ray can be pricey
- If you are asymptomatic, you wouldn't think to get an x-ray




Conclusion:

- Diagnosing infections and diseases with Computer Vision using chest x-rays work extremely well, but we need to consider the risks when giving X-rays to people who have an extremely contagious infection

How Else Can We Determine if Someone Currently Has COVID-19?




Classifying COVID cough audio via Mel-Spectrograms



$$\times \quad 2048 \quad =$$

1. Gathered cough audio from University of Stanford's Virufy dataset via Github
2. Two Classes: healthy and COVID
3. Created mel-spectrogram images for each segmented mp3 audio file
4. Trained Sequential Convolutional Neural Network model off the images



Virufy Model Evaluation Results




- 85% Accuracy
- 100% Recall for COVID class
- 0% False Negative rate

Things to note:

- The virufy dataset consisted of 121 segmented audio files from 16 different people
- May be missing important variables that could help further differentiate between classes

RECOMMENDATION:

Using cough audio to diagnose COVID-19 in infected individuals would be extremely safe, efficient, cheap, and provide virtually no risk




Combined Audio Datasets

1. Gathered CoughVid cough audio via Zenodo
2. Combined with Virufy audio dataset
3. Created mel-spectrograms for each audio file
4. Trained Sequential Neural Network model off the images

Did not achieve the accuracy and COVID class recall rate we were aiming for


Things to consider:

- Heavy class imbalance
- A lot of Silence in some of the audio
- Labels for CoughVid data were self-reported
- Model complexity



Recommendation:

Gathering high quality coughing audio from current COVID-19 patients through Hospitals and Institutions could significantly boost the model's recall rate for the COVID class



Final Conclusions

- Using Computer Vision for diagnosing diseases and infections through x-rays and spectrograms can be extremely beneficial for patients and staff (depending on the disease/infection)
- Gathering high quality COVID positive cough audio data and storing in a database would help further improve our models accuracy and recall rate, allowing us to move quicker towards the end goal of creating an app to quickly diagnose people who have COVID-19
- Setting up the process of audio data collection from various sources could help protect against future variants or strains of COVID-19 or similar diseases




Further Research:

1. Gather more higher quality data - request access from other universities or institutions that are also gathering coughing audio to diagnose COVID-19
2. feature extraction of audio files
 - a. Segmenting audio files (5 second segments)
 - b. Tweaking spectrogram parameters (sample rate, vmin, vmax, etc.)
3. Apply use of transfer models (e.g. VGG-16, ImageNet, etc.)


Thanks for your time!

Are there any questions?



References

1. Virginia Department of Health. "COVID-19 Testing." *Vdh.Virginia.Gov*, Virginia Department of Health, 5 Feb. 2021, www.vdh.virginia.gov/coronavirus/covid-19-testing/#:~:text=There%20are%20two%20main%20types,)%)20and%20antibody%20tests.
2. Shmerling, Robert H. "Which Test Is Best for COVID-19?" *Harvard Health Blog*, 5 Jan. 2021, www.health.harvard.edu/blog/which-test-is-best-for-covid-19-2020081020734#:~:text=The%20reported%20rate%20of%20false,infection%20the%20test%20is%20performed.
3. "Coronavirus Disease 2019 (COVID-19)." *Centers for Disease Control and Prevention*, 11 Feb. 2020, www.cdc.gov/coronavirus/2019-ncov/more/scientific-brief-sars-cov-2.html#:~:text=%E2%80%A2Contact%20transmission%20is%20infection,(typically%20hours).
4. Smyth, Tamara. "The Mel Scale." *University of California, San Diego*, Tamara Smyth, 4 June 2019, musicweb.ucsd.edu/~trsmyth/pitch2/Mel_Scale.html.
5. M.E.H. Chowdhury, T. Rahman, A. Khandakar, R. Mazhar, M.A. Kadir, Z.B. Mahbub, K.R. Islam, M.S. Khan, A. Iqbal, N. Al-Emadi, M.B.I. Reaz, M. T. Islam, "Can AI help in screening Viral and COVID-19 pneumonia?" *IEEE Access*, Vol. 8, 2020, pp. 132665 - 132676.
6. Lara Orlandic, Tomas Teijeiro, & David Atienza. (2020). The COUGHVID crowdsourcing dataset: A corpus for the study of large-scale cough analysis algorithms (Version 1.0) [Data set]. Zenodo. <http://doi.org/10.5281/zenodo.4048312>
7. Chaudhari, Gunvant, et al. "Virufy: Global Applicability of Crowdsourced and Clinical Datasets for AI Detection of COVID-19 from Cough." *arXiv preprint arXiv:2011.13320* (2020).
8. Hosseini M, Kooraki S, Gholamrezanezhad A, Reddy S, Myers L. Radiology Perspective of Coronavirus Disease 2019 (COVID-19): Lessons From Severe Acute Respiratory Syndrome and Middle East Respiratory Syndrome. *AJR Am J Roentgenol* 2020;214:1078-82. doi:10.2214/AJR.20.22969 pmid:32108495
9. Zafra, Miguel Fernández. "Understanding Convolutions and Pooling in Neural Networks: A Simple Explanation." *Medium*, 25 May 2020, towardsdatascience.com/understanding-convolutions-and-pooling-in-neural-networks-a-simple-explanation-885a2d78f211.



Appendix

Viral Testing

Molecular (PCR) Test:


- Detects the virus's genetic material via RT-PCR technique
- Sample is collected with a nasal or throat swab
- PCR tests (in general) are highly accurate
- Takes days to over a week to get the results

Antigen Test:

- Detects specific proteins on the coronavirus
- Referred to as rapid-diagnostic test: takes < 1 hour for results
- If you test positive you are very likely to be infected
- Higher chance of false negatives

Antibody Testing

- Uses a blood sample to look for antibodies developed as a result from past infection.
- Takes 1 - 3 weeks for a person to develop antibodies
- Having an antibody test too early can lead to false negative results
- Even if you've been exposed to the virus, it's not yet known:
 - If you're immune
 - How long immunity lasts
 - Lack of contagiousness




Our Social Defense



- Social distancing
- Use of masks in the community
- Hand hygiene
- Surface cleaning and disinfection
- Ventilation
- Avoidance of crowded indoor spaces



Age of Persons in CoughVid Dataset



CLAHE



- Contrast Limited Adaptive Histogram Equalization
- focuses on small portions of the image
- **Bilinear interpolation** - a resampling method that uses the distance-weighted average of the X nearest pixel values to estimate a new pixel value.
- enhances the contrast of the image
- helps with the pixel intensity distribution


RECOMMENDATION:

Use CLAHE as a preprocessing technique if working with x-ray or MRI scans. CLAHE is able to provide enough contrast to the image without over-amplifying the intensity of the pixels, providing more 'depth' within each image. It is a great tool if the goal of your project involves detection and/or recognition between classes.

Convolution and Pooling


Convolutional Layers:

- Applying a filter (green) to an image (blue)
- Creating a new image (red)
- Allows us to retain the important features of the image while ignoring the rest





Max Pooling Layers:

- Reduces the information in the image
- Keeps (and intensifies) the important features
- These important features are kept throughout the model





Fast Fourier Transformation




- Our waveplot (top figure) is our “signal”
- The FFT is an efficient algorithm for computing the Discrete Fourier Transform.
- The DFT (bottom figure) allows us to decompose a signal into individual frequencies and the magnitude of each frequency
- The DFT represents a histogram of the magnitude of frequencies in a sound
- No time associated with a DFT; it is static

Fast Fourier Transformation



- Our waveplot (top figure) is our “signal”
- The FFT is an efficient algorithm for computing the Discrete Fourier Transform.
- The DFT (bottom figure) allows us to decompose a signal into individual frequencies and the magnitude of each frequency
- The DFT represents a histogram of the magnitude of frequencies in a sound
- No time associated with a DFT; it is static

Spectrograms



- Each DFT is calculated on overlapping windowed portions of the signal
- A spectrogram is basically composed of multiple Discrete Fourier Transforms
- In order for us to visualize “loudness” in our signal, we must convert from amplitude to decibels
- Allows us to view the loudness of frequencies over time

Mel-Spectrograms

- A spectrogram where the frequencies are converted to the mel-scale
- The **mel-scale** = “a perceptual scale of pitches judged by listeners to be equal in distance from one another”
- Perceptually relevant amplitude and frequency representation

